

Toronto Ferry Ticket Demand Forecasting using SARIMA and Prophet with MLOps Integration (DVC & MLflow)

This project demonstrates a complete time series forecasting workflow using SARIMA and Prophet. It includes preprocessing, exploratory analysis, model fitting, and forecasting for a 30-day horizon. The pipeline is modular, reproducible, and designed for extensibility.

Project structure

```
sarima-prophet-mlflow-dvc-pipeline/
├── .dvc/                                # DVC configuration
├── .venv/                                # Virtual environment
├── data/
│   ├── raw/
│   │   └── TorontoIslandFerryTicketCount.csv
│   └── processed/
│       └── daily_tickets_2022_25.parquet
├── models/
│   ├── sarima_model.pkl
│   └── prophet_model.pkl
├── notebooks/
│   ├── exploratory_analysis.ipynb
│   ├── sarima_exploration.ipynb
│   └── prophet_exploration.ipynb
├── output/
│   ├── plots/
│   └── results/
│       ├── sarima_cv_results.csv
│       └── prophet_cv_results.csv
├── scripts/
│   ├── __init__.py
│   ├── sarima_model.py
│   └── prophet_model.py
├── utils/
│   ├── __init__.py
│   ├── evaluate_forecast.py
│   └── mlflow_logger.py
└── .env
└── requirements.txt
└── README.md
```

Toronto Ferry Ticket Sales Redemption Data





Figure: Daily redemption count from 2022 onwards.

- Yearly seasonality - winter low and summer high
- Weekly seasonality - weekend high
- Presence of spikes on summer days

Basic Statistics

- The mean ticket count is approximately 3,700, while the median (50th percentile) is around 1,600, indicating a right-skewed distribution. This suggests that a few extreme high values are pulling the mean above the median.
- A high standard deviation of 4,440 tickets reflects a wide dispersion in daily ticket sales, pointing to substantial variability across days.

Log-transformation for variance stability and first-difference to remove trend and make the series stationary



*Figure: log-transformed daily redemption stabilizes the variance.

First differencing of log-transformed series





Figure: We see that the series mean is zero, with almost stable variance and sudden spikes

One-step differencing of log-transformed time-series

- Log transformation stabilizes the variance and reduce the impact of spikes

- one-step differencing now removes the trend, making the time-series stationary.
- one-step differencing represents the relative changes or growth rate in redemption counts from one day to the next.
- it centers around zero with occasional spikes and dips

ACF and PACF plots



Auto-correlation and Partial Auto-correlation of 1-step differencing of log-transformed time series

- PACF suggests significant auto-correlation at lag 1, 2, 3, 4, 5 and 6. After that the correlation dies. We can try with AR(p) p from 1 to 6 and see the BIC
- ACF suggests significant relationships at lag 1 and 2. q could be 1 or 12

Augmented Dickey-Fuller test

```
from statsmodels.tsa.stattools import adfuller
import pandas as pd

# Run the test
result = adfuller(data_diff['diff_redemption_count'].dropna())

# Organize results
adf_output = {
    'ADF Statistic': result[0],
    'p-value': result[1],
    'Number of Lags Used': result[2],
    'Number of Observations Used': result[3],
    'Critical Values': result[4],
    'IC Best (AIC)': result[5]
}

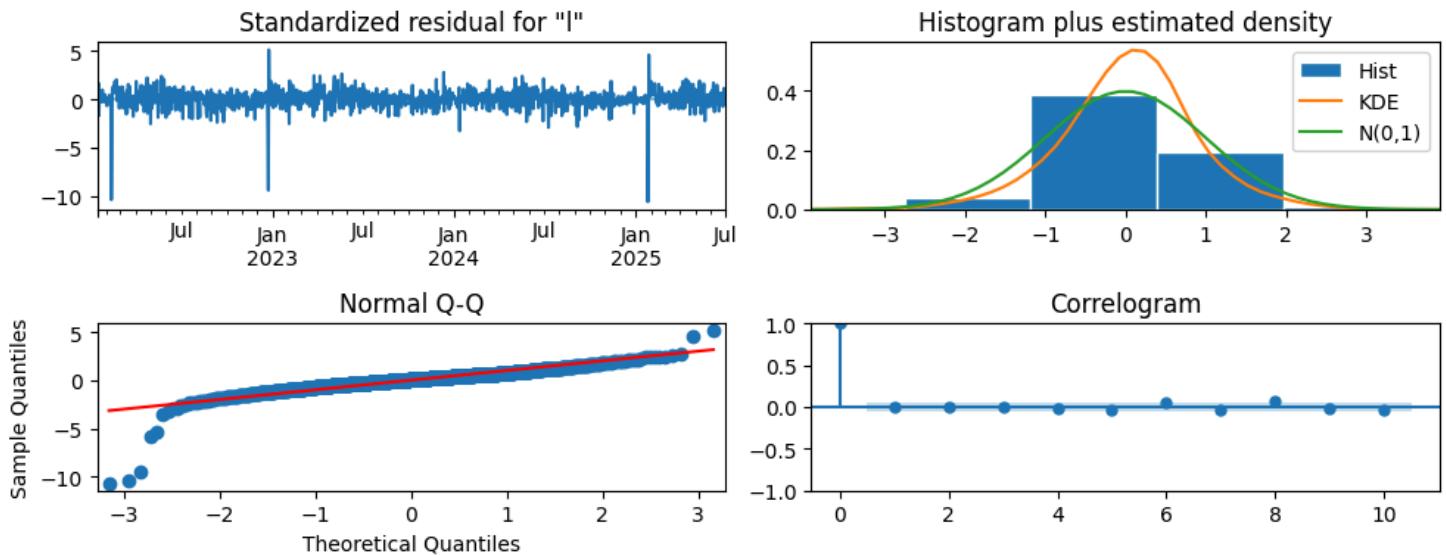
# Display as DataFrame
adf_df = pd.DataFrame.from_dict(adf_output, orient='index', columns=['Value'])
print(adf_df)
```

	Value
ADF Statistic	-16.285116
p-value	0.0
Number of Lags Used	12
Number of Observations Used	1264
Critical Values	{'1%': -3.4355, '5%': -2.8638}
IC Best (AIC)	2274.831503

- p-value = 0 suggests that the data does not provide enough evidence in favor of the null hypothesis that the time series has a unit root - the series is non-stationary.
- So we reject the Null Hypothesis that the series is non-stationary.
- 1-step differencing or log-transformation has made the series stationary.

Apply SARIMA (2,1,2) \times (1,0,1,7) model on log-transformed redemption tickets

```
sarima_model = SARIMAX(  
    y,  
    order=(2,1,2),  
    seasonal_order=(1,0,1,7),  
    enforce_stationarity=False,  
    enforce_invertibility=False  
)  
final_res = sarima_model.fit(disp=False)  
  
forecast_steps = 30 # 7  
future_index = pd.date_range(start=y.index[-1] + pd.Timedelta(days=1), periods=forecast_steps, freq='D')  
forecast = final_res.forecast(steps=forecast_steps)  
forecasted_object = final_res.get_forecast(steps=forecast_steps)  
mean_forecast = forecasted_object.predicted_mean  
confint = forecasted_object.conf_int()  
  
forecast = pd.Series(forecast, index=future_index)
```



- Residuals are approximately normal, but not perfect. Mild departure from normality, especially in the extremes. This is acceptable for forecasting, though you might consider a robustness check if prediction intervals are critical.
- The spike at lag 1 (Correlogram) suggests some autocorrelation remains, which means the model might be slightly underfitting short-term dynamics.
- SARIMA(2,1,2)(1,0,1,7) model is reasonably well-specified.

Prophet Modeling with TimeSeriesSplit Cross-Validation

```
# TimeSeriesSplit with 7-day test windows to simulate weekly forecasting
tscv = TimeSeriesSplit(n_splits=n_splits, test_size=forecast_horizon)

for fold, (train_idx, test_idx) in enumerate(tscv.split(df)):


    train_df = df.iloc[train_idx]
    test_df = df.iloc[test_idx]

    model = Prophet(
        growth='linear',
        yearly_seasonality=True,
        weekly_seasonality=True,
        daily_seasonality=False,
        holidays=holiday_df
    )
    model.fit(train_df)


    future = model.make_future_dataframe(periods=len(test_df), freq='D')
    forecast = model.predict(future)
    # forecast_test = forecast.set_index('ds').loc[test_df['ds'].values]
    forecast_test = forecast.set_index('ds').reindex(test_df['ds'])

    evaluation_metric_dict = evaluate_forecast(test_df['y'], forecast_test['yhat'])
```

Fold 5: Prophet Prediction (log-transformed)



Fold 5: Prophet Prediction (Original Scale)



Model Evaluation: SARIMA vs Prophet (30-day forecast)

A comparison table below showing **SARIMA vs Prophet** on original scale metrics, with **% improvement** of Prophet over SARIMA for each fold:

Fold-wise Comparison with % Improvement

Fold	SARIM A MAE	Prophet MAE	% MAE ↓	SARIM A RMSE	Prophet RMSE	% RMSE ↓	SARIM A SMAPE	Prophet SMAPE	% SMAPE ↓
1	473.04	213.04	54.98%	507.79	246.57	51.44%	101.30	33.55	66.88%
2	265.57	267.11	-0.58%	322.59	371.62	-15.20%	31.37	30.14	3.89%
3	1273.46	819.95	35.61%	1770.19	1164.53	34.23%	72.41	41.31	42.94%
4	1984.88	1531.90	22.78%	3091.61	2152.43	30.37%	50.77	41.67	17.91%
5	4368.85	3167.34	27.47%	5954.60	4324.06	27.42%	59.20	39.96	32.52%

$$\% \text{ Improvement} = ((\text{SARIMA} - \text{Prophet}) / \text{SARIMA}) \times 100$$

Aggregated Summary (Mean across folds)

Metric	SARIMA	Prophet	% Improvement
MAE	1873.16	1199.47	35.95%
RMSE	2329.76	1251.85	46.27%
SMAPE	62.81	37.33	40.58%

Insights

- Prophet consistently outperforms SARIMA in Fold 1, 3, 4, and 5, with substantial gains in SMAPE and RMSE.
- Fold 2 is an anomaly where SARIMA slightly edges Prophet in MAE and RMSE.
- Prophet achieves ~36% lower MAE, ~46% lower RMSE, and ~41% lower SMAPE on average. If you're selecting a model for forecasting, Prophet appears to generalize better across varied error metrics.
- Overall, **Prophet** shows **stronger generalization** capturing both trend and seasonality and **lower error** across all metrics.

MLflow Integration with DagsHub

This project uses MLflow for experiment tracking, model management, and reproducibility, fully integrated with DagsHub's hosted MLflow server.

dagshub.com/elias.reaz/dvc-mlflow-forecasting-pipeline/experiments#experiment/m_af685cc764c0432a8fc0a2a6b69c11bf

Google Chrome isn't your default browser Set as default

Search public repositories Issues Pull Requests Resources Explore Pricing Create +

elias.reaz / dvc-mlflow-forecasting-pipeline

Experiments 23 Models Annotations Collaboration Settings

Unwatch 1 Star Fork

Register Model Delete

Experiments < SARIMA_Modeling

Details

Parameters

Status: Finished

Commit ID: cf4e07a3ee

Source: (copy)

Data: (copy)

Duration: 1min

Date: Nov 12, 2025 at 0:06:05

Labels: +

MLflow run ID: af685cc764c0432a8fc0a2a6b69c11bf

Parameters

Name	Value
model	tree
model_type	SARIMA
sarima_order	[2, 1, 2]
sarima_seasonal_order	[1, 0, 1, 7]

Metrics

Name	Value
fold_1_mae_log	1.1584
fold_1_mae_original	473.0444
fold_1_rmse_log	1.2343
fold_1_rmse_original	507.7946
fold_1_smape_log	19.6192

<https://dagshub.com/elias.reaz/dvc-mlflow-forecasting-pipeline/experiments>